


Reasoning: Load the json file into a pandas DataFrame.

```
import pandas as pd
try:
    df_hotels = pd.read_json('filtered_hotels.json').head(2)
    display(df_hotels.head())
    display(df_hotels.info())
except FileNotFoundError:
    print("Error: 'filtered_hotels.json' not found.")
    df_hotels = None
except ValueError as e:
    print(f"Error: Could not parse JSON file: {e}")
    df_hotels = None
except Exception as e:
    print(f"An unexpected error occurred: {e}")
    df_hotels = None
```



	hotel_id	name	location	price	rating	features	nearby_beaches	comments
0	1	sousse pearl marriott resort spa	Sousse	152	Avec une note de 8,2	['4 piscines', 'Navette aéroport', 'Chambres n...	['Plage de Boujaafar\n500 m', 'Plage de Bhar E...	['titre': 'Exceptionnel', 'commentaire': 'J'a...
1	2	hotel medina	Sousse	75	Avec une note de 7,7	['Chambres non-fumeurs', 'Parking', 'Connexion...	['Plage de Bhar Ezzebla\n750 m', 'Plage de Bou...	['titre': 'Exceptionnel', 'commentaire': 'Bel...

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2 entries, 0 to 1
Data columns (total 8 columns):
Column Non-Null Count Dtype
--- ----- -
0 hotel_id 2 non-null int64
1 name 2 non-null object
2 location 2 non-null object
3 price 2 non-null int64
4 rating 2 non-null object
5 features 2 non-null object
6 nearby_beaches 2 non-null object
7 comments 2 non-null object
dtypes: int64(2), object(6)
memory usage: 760 B+ bytes


Data preparation

Subtask:

Prepare the data for embedding generation.

▼ Note wl Prix yekhdmo

```
if df_hotels is not None and 'nearby_beaches' in df_hotels.columns:
    for index, row in df_hotels.iterrows():
        if isinstance(row['nearby_beaches'], str):
            df_hotels.at[index, 'nearby_beaches'] = row['nearby_beaches'].replace('\n', ' ')
            print(df_hotels["nearby_beaches"])
        else:
            print("DataFrame or 'nearby_beaches' column not found.")
df_hotels["nearby_beaches"]
```



```
0   ['Plage de Boujaafar 500 m', 'Plage de Bhar Ez...
1   ['Plage de Bhar Ezzebla 750 m', 'Plage de Bouj...
Name: nearby_beaches, dtype: object
```

	nearby_beaches
0	['Plage de Boujaafar 500 m', 'Plage de Bhar Ez...
1	['Plage de Bhar Ezzebla 750 m', 'Plage de Bouj...

dtype: object

```

# df_hotels['text'] = (
#     df_hotels['name'].fillna('Nom inconnu') + ', ' +
#     df_hotels['location'].fillna('Localisation inconnue') + ', ' +
#     df_hotels['price'].fillna('Prix non renseigné').astype(str) + ' DT, ' +
#     df_hotels['rating'].fillna('Note non disponible').astype(str)
# )

# # Inspect the 'text' column
# display(df_hotels.head())
# display(df_hotels['text'].unique())

import ast
# clean the features and beaches
def clean_list_field(field_value):
    try:
        items = ast.literal_eval(field_value)
        if isinstance(items, list):
            cleaned = list(dict.fromkeys([item.replace('\n', ' ').strip().lower() for item in items if isinstance(item, str)]))
            return ', '.join(cleaned)
        except Exception:
            pass
        return 'non renseigné'
# hezz ken l comments
def extract_individual_comments(row):
    """
    Extrait et nettoie les commentaires dans une liste séparée,
    supprimant les '\n' et ne gardant que le contenu du commentaire.
    """
    comments = row['comments']
    if isinstance(comments, list):
        # Nettoyer chaque commentaire et enlever les '\n' (saute de ligne)
        cleaned_comments = [c.get('commentaire', '').replace("\n", " ") for c in comments]
        return cleaned_comments
    else:
        return ['Aucun commentaire disponible']

# Appliquer cette fonction à chaque ligne du DataFrame pour extraire les commentaires
df_hotels['individual_comments'] = df_hotels.apply(extract_individual_comments, axis=1)

def generate_text(row):
    name = row.get('name', 'nom inconnu')
    location = row.get('location', 'localisation inconnue')
    price = row.get('price', 'prix non renseigné')
    rating = row.get('rating', 'note non disponible')

    features = clean_list_field(row.get('features', ''))
    beaches = clean_list_field(row.get('nearby_beaches', ''))

    # les commentaires f liste whdhom donc bch naamloulhom jointure
    # Traitement des commentaires
    comments = row['individual_comments']
    if comments:
        # Afficher tous les commentaires individuellement
        cleaned_comments = ' | '.join([f"{c}" for c in comments])
    else:
        cleaned_comments = 'Aucun commentaire disponible'

    return (
        f"Nom de l'hôtel : {name}.\n"
        f"Localisation : {location}.\n"
        f"Prix : {price} TND par nuit.\n"
        f>Note : {rating}.\n"
        f"Caractéristiques : {features}.\n"
        f"Plages à proximité : {beaches}.\n"
        f"Avis clients : {cleaned_comments}."
    )

df_hotels['text'] = df_hotels.apply(generate_text, axis=1)

```

✓ Feature engineering

Subtask:

Generate embeddings for the 'text' column in the `df_hotels` DataFrame using Hugging Face embeddings.

```
!pip install langchain langchain-community pypdf chromadb -q
!pip install langchain_groq -q
!pip install -U langchain-huggingface -q
!pip install -U langchain-chroma -q
!pip install gradio -q
```

```

67.3/67.3 kB 4.7 MB/s eta 0:00:00
Installing build dependencies ... done
Getting requirements to build wheel ... done
Preparing metadata (pyproject.toml) ... done
2.5/2.5 MB 52.2 MB/s eta 0:00:00
303.4/303.4 kB 24.3 MB/s eta 0:00:00
18.9/18.9 MB 38.4 MB/s eta 0:00:00
94.9/94.9 kB 9.9 MB/s eta 0:00:00
284.2/284.2 kB 26.9 MB/s eta 0:00:00
2.0/2.0 MB 88.8 MB/s eta 0:00:00
101.6/101.6 kB 10.5 MB/s eta 0:00:00
16.4/16.4 MB 114.9 MB/s eta 0:00:00
55.9/55.9 kB 5.5 MB/s eta 0:00:00
194.9/194.9 kB 20.1 MB/s eta 0:00:00
65.8/65.8 kB 6.3 MB/s eta 0:00:00
118.9/118.9 kB 12.3 MB/s eta 0:00:00
92.0/92.0 kB 10.1 MB/s eta 0:00:00
44.4/44.4 kB 4.3 MB/s eta 0:00:00
62.5/62.5 kB 6.4 MB/s eta 0:00:00
459.8/459.8 kB 41.3 MB/s eta 0:00:00
50.9/50.9 kB 5.9 MB/s eta 0:00:00
71.5/71.5 kB 7.7 MB/s eta 0:00:00
4.0/4.0 MB 102.3 MB/s eta 0:00:00
454.8/454.8 kB 39.7 MB/s eta 0:00:00
46.0/46.0 kB 606.8 kB/s eta 0:00:00
86.8/86.8 kB 8.9 MB/s eta 0:00:00
Building wheel for pypika (pyproject.toml) ... done
127.5/127.5 kB 5.8 MB/s eta 0:00:00
437.7/437.7 kB 15.4 MB/s eta 0:00:00
363.4/363.4 MB 3.9 MB/s eta 0:00:00
13.8/13.8 MB 102.4 MB/s eta 0:00:00
24.6/24.6 MB 89.1 MB/s eta 0:00:00
883.7/883.7 kB 57.1 MB/s eta 0:00:00
664.8/664.8 MB 1.3 MB/s eta 0:00:00
211.5/211.5 MB 5.7 MB/s eta 0:00:00
56.3/56.3 MB 12.9 MB/s eta 0:00:00
127.9/127.9 MB 7.5 MB/s eta 0:00:00
207.5/207.5 MB 5.5 MB/s eta 0:00:00
21.1/21.1 MB 80.0 MB/s eta 0:00:00
611.1/611.1 kB 14.4 MB/s eta 0:00:00
2.4/2.4 MB 82.2 MB/s eta 0:00:00
54.1/54.1 MB 16.3 MB/s eta 0:00:00
322.9/322.9 kB 28.9 MB/s eta 0:00:00
11.5/11.5 MB 133.7 MB/s eta 0:00:00

```

```
from langchain.text_splitter import RecursiveCharacterTextSplitter
from langchain_huggingface import HuggingFaceEmbeddings
from langchain_community.vectorstores import Chroma
from langchain.chains import RetrievalQA
from langchain_groq import ChatGroq
from langchain.prompts import PromptTemplate
from langchain.schema import Document
```

```
import os
import gradio as gr
import json
```

```
!pip install language_tool_python
```

```

Collecting language_tool_python
  Downloading language_tool_python-2.9.3-py3-none-any.whl.metadata (54 kB)
    54.7/54.7 kB 3.6 MB/s eta 0:00:00
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from language_tool_python) (2.32.3)

```

```
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from language_tool_python) (4.67.1)
Requirement already satisfied: psutil in /usr/local/lib/python3.11/dist-packages (from language_tool_python) (5.9.5)
Requirement already satisfied: toml in /usr/local/lib/python3.11/dist-packages (from language_tool_python) (0.10.2)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->language_tool_python) (3.4.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests->language_tool_python) (3.10.1)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests->language_tool_python) (2.2.3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests->language_tool_python) (2025.11.11)
Downloading language_tool_python-2.9.3-py3-none-any.whl (55 kB)
55.1/55.1 kB 4.3 MB/s eta 0:00:00
Installing collected packages: language_tool_python
Successfully installed language_tool_python-2.9.3
```


pour corriger le probleme des reponses j'ai changer le model name vers e5-large khtr yhezz akther tokens hata l 32k ama yab9a 9rabet 45 min bch ylanci

```
# Initialize the Hugging Face embeddings model
#embedding_function = HuggingFaceEmbeddings(model_name="sentence-transformers/all-MiniLM-L6-v2")

# switching to another model with more tokens
embedding_function = HuggingFaceEmbeddings(model_name="intfloat/e5-large-v2")

# Generate embeddings for the 'text' column
hotel_embeddings = embedding_function.embed_documents(df_hotels['text'].tolist())

# Print the shape of the embeddings
print(f"Shape of the embeddings: {len(hotel_embeddings)}, {len(hotel_embeddings[0])}")
```

 /usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning: The secret `HF_TOKEN` does not exist in your Colab secrets. To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>). You will be able to reuse this secret in all of your notebooks. Please note that authentication is recommended but still optional to access public models or datasets.

```
warnings.warn(
modules.json: 100% 387/387 [00:00<00:00, 37.6kB/s]
README.md: 100% 67.8k/67.8k [00:00<00:00, 5.25MB/s]
sentence_bert_config.json: 100% 57.0/57.0 [00:00<00:00, 5.04kB/s]
config.json: 100% 616/616 [00:00<00:00, 51.7kB/s]
model.safetensors: 100% 1.34G/1.34G [00:08<00:00, 188MB/s]
tokenizer_config.json: 100% 314/314 [00:00<00:00, 33.3kB/s]
vocab.txt: 100% 232k/232k [00:00<00:00, 3.34MB/s]
tokenizer.json: 100% 711k/711k [00:00<00:00, 4.97MB/s]
special_tokens_map.json: 100% 125/125 [00:00<00:00, 13.4kB/s]
config.json: 100% 201/201 [00:00<00:00, 21.5kB/s]
Shape of the embeddings: 2, 1024
```

9a3ed nfasakh fl chroma db f kol mara bch yaawd ysavi les index jdod , solution lel caractéristique heya eni zedet fl chunk size wl overlap

✓ kif zedet chunk brcha wala ykhalwedh w ya9rach f kol chy

```
import shutil
shutil.rmtree("chroma_db", ignore_errors=True)

from langchain_community.vectorstores import Chroma
from langchain.schema import Document
from langchain.text_splitter import RecursiveCharacterTextSplitter

# Split the text into chunks
splitter = RecursiveCharacterTextSplitter(chunk_size=1000, chunk_overlap=300)
documents = [Document(page_content=t) for t in df_hotels['text'].tolist()]
```

```

split_docs = splitter.split_documents(documents)

# Persisted vectorstore
persist_directory = "chroma_db"

if os.path.exists(persist_directory):
    print("Chargement de la base Chroma existante...")
    vectorstore = Chroma(persist_directory=persist_directory, embedding_function=embedding_function)
else:
    print("Création d'une nouvelle base Chroma...")
    vectorstore = Chroma.from_documents(
        documents=split_docs,
        embedding=embedding_function,
        persist_directory=persist_directory,
        collection_name="hotels",
    )
vectorstore.persist()

# Now reload the persisted vectorstore
vectorstore = Chroma(
    embedding_function=embedding_function,
    persist_directory=persist_directory,
    collection_name="hotels"
)

```

Création d'une nouvelle base Chroma...

<ipython-input-14-23a6dcc2309f>:25: LangChainDeprecationWarning: Since Chroma 0.4.x the manual persistence method is r
vectorstore.persist()
<ipython-input-14-23a6dcc2309f>:28: LangChainDeprecationWarning: The class `Chroma` was deprecated in LangChain 0.2.9
vectorstore = Chroma(

✓ test fi chroma DB

```

print(f"Nombre de documents indexés : {vectorstore._collection.count()}")
for doc in split_docs:
    if "medina" in doc.page_content.lower():
        print(doc.page_content)

```

Nombre de documents indexés : 7
 Nom de l'hôtel : hotel medina.
 Localisation : Sousse.
 Prix : 75 TND par nuit.
 Note : Avec une note de 7,7.
 Caractéristiques : chambres non-fumeurs, parking, connexion wi-fi gratuite, chambres familiales, bar, très bon petit-c
 Plages à proximité : plage de bhar ezzebla 750 m, plage de boujaafar 1,1 km, las vegas beach 4,1 km, plage du thalass
 Avis clients : Bel hôtel proche plage et medina | L'emplacement est impeccable. La chambre était propore. Le petit déj
 L' accueil est familial et chaleureux. Le petit déjeuner est excellent. A peine sorti, on est à la Grande Mosquée et a
 client et très prévenante. Nous avons aimé notre séjour dans cet hôtel. | Hôtel à l'entrée de la médina. Chambre famil

```

# Initialize the ChatGroq LLM
llm = ChatGroq(model="llama-3.3-70b-versatile", api_key="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")

# Create a RetrievalQA chain#

#qa_chain = RetrievalQA.from_chain_type(
#    llm=llm, retriever=vectorstore.as_retriever()
#)

```

j'ai testé le mixtral-8x7b mais il donne un erreur

```

llm = ChatGroq(model="llama-3.3-70b-versatile", api_key="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")

prompt_template = PromptTemplate(
    template=(
        "Tu es un assistant touristique tunisien spécialisé dans les recommandations d'hôtels.\n"
        "Utilise uniquement les informations fournies dans le contexte ci-dessous.\n"
    )
)

```

```
    "**N'invente jamais.** Si une information est absente, indique : 'Non renseigné'.\n\n"
    "Contexte : {context}\n"
    "Question du client : {question}\n\n"
    "Réponse détaillée:"
),
input_variables=["context", "question"]
)

qa_chain = RetrievalQA.from_chain_type(
    llm=llm,
    chain_type="stuff",
    retriever=vectorstore.as_retriever(),
    return_source_documents=True,
    chain_type_kwargs={"prompt": prompt_template}
)

def chatbot(query):
    result = qa_chain({"query": query})
    return result['result']

iface = gr.Interface(
    fn=chatbot,
    inputs=gr.Textbox(lines=2, placeholder="Enter your question here..."),
    outputs="text",
    title="Hotel RAG Chatbot",
)

iface.launch()
```

🔗 It looks like you are running Gradio on a hosted a Jupyter notebook. For the Gradio app to work, sharing must be enabled.

Colab notebook detected. To show errors in colab notebook, set debug=True in launch()

* Running on public URL: <https://2e35029f4f2aa1e23a.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal

query

tu connais hotel medina

Effacer

Envoyer

output

Chambres non fumeuses

- * Parking
- * Connexion Wi-Fi gratuite
- * Chambres familiales
- * Bar
- * Petit-déjeuner inclus, très copieux et apprécié par les clients

L'hôtel est situé à proximité de plusieurs plages, notamment :

- * Plage de Bhar Ezzebla (750 m)
- * Plage de Boujaafar (1,1 km)
- * Las Vegas Beach (4,1 km)
- * Plage du Thalassa Sousse (5 km)
- * Plage Hammam Sousse (8 km)

Les clients ont généralement apprécié leur séjour à l'hôtel Medina, soulignant la qualité de l'accueil, la décoration, la propreté et la situation de l'établissement. Ils ont également apprécié le petit-déjeuner, qui est décrit comme excellent et copieux. L'emplacement de l'hôtel est également

